



## AI Engineer

Location

San Francisco, New York

Type

Full time

Department

Engineering

### OverviewApplication

Large Language Models (LLMs) represent one of the most important technical advancements of the century. However, empowering our customers to fully harness these requires not just raw LLM capabilities but also skilled individuals who can build applications on top of them. The creation of ChatGPT wouldn't have been possible with only ML engineers. We are assembling a diverse team of ML, Frontend, Software engineers, Designers, and AI Strategists, to drive the adoption of AI powered solutions across the F500.

#### Responsibilities:

- Hit the ground running and take ownership of strategic accounts with iconic customers.
- Thoughtfully design and implement systems that leverage LLMs in a responsible and predictable manner. This includes model vetting, pilot development, and production implementation.
- Collaborate with customers to understand their challenges and pain-points, and create innovative yet practical AI solutions.
- Contribute to the development of Distyl's internal platform (Distillery) for LLM application development, including designing and implementing new features and functionality.
- Monitor and maintain production systems to ensure they are meeting system requirements and delivering the expected results

- Stay up-to-date with AI and ML trends, and leverage your practical experience with Large Language Models (especially OpenAI's) to responsibly evaluate their trade-offs and bring innovative solutions to our customers.
- Continuously improve solution design and implementation processes, incorporating feedback and lessons learned from previous projects.

## Requirements

- Bachelor's degree in Computer Science, Mathematics, or related fields.
- At least 3 years of experience as a software engineer or machine learning engineer.
- Practical experience in software development (ideally full stack) and production deployments for data-intensive or ML-based products.
- An obsession with thorough AI/ML evaluation to create reliable applications.
- Strong knowledge of machine learning fundamentals and deep learning frameworks (e.g., PyTorch, Transformers, Scikit-learn, NumPy, Pandas).
- Work experience with software development best practices, including version control (e.g., Git), automated testing, and continuous integration/continuous deployment (CI/CD).
- Experience architecting and implementing cloud data infrastructure and pipelines (Azure, AWS, and GCP).
- Strong communication and collaboration skills.

## What we offer

- Opportunity to build a mission-oriented company together to create practical adoption of the biggest revolution in human productivity.
- Opportunity to own \$Xm+ accounts with some of the most iconic institutions, and drive transformative outcomes.
- Opportunity to work on cutting-edge AI technologies and make a real impact on the adoption of LLMs in the enterprise space.
- Competitive salary and benefits package, including equity options and healthcare.
- Collaborative and supportive work environment, with a strong focus on personal and professional growth.
- Exposure to a diverse range of real-world industries and projects, providing unique challenges and learning opportunities.
- Chance to be part of a growing team that is passionate about AI and its potential to revolutionize human productivity.

If you are an innovative, ambitious, and driven individual looking to make a difference in the world of AI, we want to hear from you. Apply now to join our team as an Applied AI Engineer and help us shape the future of AI-driven solutions for enterprises across the globe.

***Note: Distyl is a hybrid working environment and requires in office collaboration 3 days a week. We have offices in SF and NYC***

Apply for this Job

This site is protected by reCAPTCHA and the Google [Privacy Policy](#) and [Terms of Service](#) apply.