AI Product Engineer

Location

San Francisco

Type

Full time

Department

Engineering

Compensation

$170K – $280K • Offers Equity

At Speak, our compensation philosophy is based on work experience, interview performance, leveling and other factors evaluated during the hiring process.

Overview Application

About us

Our mission is to reinvent the way people learn, starting with language. We begin by teaching the next billion people English, Spanish, and French.

English is the global language of business, culture, and communication, and over 1.5 billion people around the world are actively trying to learn right now. Others dream of communicating with the half-billion native Spanish speakers across the globe. The problem is that it's nearly impossible to learn to *speak* a language without constant access to a speaking partner. Grammar and vocab apps don't really help – you need to actually *converse* with someone.

Speak is on a journey to fix this. We're creating an AI-powered experience that replicates the flow of a conversation, *without* needing a human on the other end. The goal is to make it radically more accessible to be able to have conversations in a foreign language and eventually help hundreds of millions of people gain fluency who otherwise wouldn't be able to.

We started on this journey over five years ago and we've still got a long ways to go. We're thoughtfully adding new team members only when we think they can truly play a big role in our mission.

Speak launched first in South Korea where we have quickly grown to become the top grossing education app in the country. We have now delivered this winning product to more than 30 countries globally and are continuing to expand to more markets in the coming months. The company is well funded, raising a recent Series B backed by investors like OpenAI, Founders Fund, Y Combinator, Khosla Ventures, Lachy Groom, Josh Buckley, and others. We're a team of 75 based primarily in SF, Seoul, Tokyo, and Ljubljana.

About this role

As an AI Product Engineer at Speak, you'll play a pivotal role in developing the future of language learning and creating the most effective path to language fluency. Your primary responsibility will be developing innovative product experiences powered by Language Learning Models (LLMs) and deploying them to millions of users.

Your tasks will span the full stack and beyond, from evaluating models and testing prompts to developing and iterating on various product features such as our conversational onboarding and lesson experiences, grammatical assessment and feedback, and personalized lessons.

We're constantly pushing the boundaries of what LLMs can do to provide an exceptional and unparalleled language learning experience to users in over 30 countries worldwide.

What you'll be doing

- Developing and deploying LLM-powered language learning products across the full stack, as well as enhancing the quality and performance of existing AI-powered features within Speak

- Collaborating cross-functionally with other Engineering teams, Applied ML, Product, Design, and Content

- Refining our process for building LLM apps, including best practices for prompting, experimentation/evaluation, LLM Ops, measuring quality and performance, etc.

- Scaling existing product features to many more users and languages

What we're looking for

- 3-5+ years of experience in full stack/backend, product-focused software engineering

    - Proficiency in React/Node/Typescript and Python

- Strong product intuition — the ability to think broadly and cross-functionally about innovative LLM-powered capabilities and product experiences

- Ability to work independently and build at a high velocity

- Interest in building with LLMs

- You may have participated in LLM hackathons, developed several side projects, experimented with different types of models, and use ChatGPT daily

Bonus:link

- You have real-world experience developing and deploying LLM apps and a strong understanding, gained through experience, of what works and what doesn't

- A keen intuition for improving performance and output quality of LLM systems

- Experience with LLM Ops and tools (e.g., vector databases, RAG, prompt ops)

Office

- San Francisco, CA

Why work at Speak

1. **Join a fantastic, tight-knit team at the right time:** we're growing very quickly, we've raised our Series B and an additional extension from some of the top investors in the valley, and we've achieved product-market fit in our initial markets. You'd join at a magical time when a single person could significantly change the course of the company.

2. **Do your life's work with people you'll love working with:** we care strongly about our craft and want every person at Speak to feel like they're growing every day. We believe in the idea that working with people you both enjoy and have respect for makes everything better. We hire thoughtfully and only work with people we admire deeply.

3. **Global in nature:** We're live in over 40 countries and launching in a number of new markets soon. We have dedicated offices in San Francisco, Ljubljana, Seoul, and Tokyo, and you'll have the opportunity to talk to users in each of these regions on a regular basis as well as travel.

4. **Impact people's lives in a major way:** Learning a language is one of the single most life-changing skills one can learn, and right now 99% of people never achieve their goal because the process is broken. We're helping millions of people achieve their goals and improve their lives.

Speak does not discriminate based upon race, religion, color, national origin, gender (including pregnancy, childbirth, or related medical conditions), sexual orientation, gender identity, gender expression, age, status as a protected veteran, status as an individual with a disability, or other applicable legally protected characteristics.

Apply for this Job